

Minseob Kim

☎ 010-7918-2385 | ✉ caley2380@kaist.ac.kr | [GitHub](#) | [LinkedIn](#)

OBJECTIVE

My long-term goal is to build general-purpose AI agents that can reason, plan, and make reliable decisions over long horizons. I am especially interested in decision-making agents, offline reinforcement learning, and reinforcement learning from human feedback for training capable and aligned AI systems.

- **Decision-Making Agents:** Learning agents that can plan, act, and adapt across complex multi-step tasks.
- **Offline Reinforcement Learning:** Learning effective policies from previously collected trajectories, with an emphasis on long-horizon credit assignment and distribution shift.
- **Reinforcement Learning from Human Feedback:** Using preference feedback and reward modeling to improve agent behavior beyond supervised learning.
- **Reliable AI Agents:** Understanding and mitigating reward hacking, unintended misalignment, and failure modes in autonomous systems.

EDUCATION

Korea Advanced Institute of Science and Technology (KAIST)

B.S. in School of Computing

Mar. 2023 – Feb. 2029 (expected)

Daejeon, South Korea

RELEVANT COURSEWORK

AI/ML Foundations: Probability and Statistics (A+), Introduction to Algorithms (A-), Data Structure (A-), Machine Learning ([self-studied](#))

Programming and Systems: Introduction to Programming (A+), System Programming (A)

PROJECTS

Paper Pilot

[\[GitHub\]](#)

AI-powered research paper reading assistant

RAG, Grounded QA

- Built a local-first research-paper reader with RAG-backed Q&A, page-aligned translation, figure/region explanations, highlights, summaries, citation cards, and persistent paper memory.
- Implemented local retrieval over extracted PDF text using overlapping chunks and lexical ranking, grounding agent answers in retrieved page excerpts with inline page citations.

Locium

[\[GitHub\]](#)

Structured memory workspace for AI agents

Agent Memory, Knowledge Representation

- Built a local-first memory workspace that converts notes, documents, transcripts, and imported text into a browsable wiki and an AI-queryable structured memory graph.
- Designed a provenance-aware memory architecture with immutable raw sources, append-only events, wiki/graph/search projections, and structured recall for AI agents.

SKILLS

Programming: C, C++, Python, PyTorch

Languages: Korean (Native), English: Professional proficiency (TOEFL iBT 5.5/6.0, Mar. 2026)